

COMPUTING RISK!!!

Question 1

In [59]: `from IPython.display import Image`
 Image("/content/sample_data/Screen Shot 2020-02-22 at 9.39.15 AM.png")

Out[59]:

Computing Risk.

1. $\mathbf{x} \in \mathbb{R}^n$ $\{-2, -1, 0, 1, 2\}$ $P = 1/5$

(a) $E[\|\mathbf{x}\|_2^2]$

$$\rightarrow E[x_1^2 + x_2^2 + \dots + x_n^2]$$

$$\rightarrow E[x_1^2] + E[x_2^2] + \dots + E[x_n^2]$$

$$\rightarrow n E[x_1^2]$$

$$\rightarrow n \left[(-2)^2 \cdot \frac{1}{5} + (-1)^2 \cdot \frac{1}{5} + 0^2 \cdot \frac{1}{5} + 1^2 \cdot \frac{1}{5} + 2^2 \cdot \frac{1}{5} \right]$$

$$\rightarrow \underline{2n}$$

(b) $E[\|\mathbf{x}\|_{\max}]$

$$\rightarrow E[\max(|x_1|, |x_2|, \dots, |x_n|)]$$

$$\rightarrow \sum_{i=0}^n i \cdot P(\max(|x_1|, \dots, |x_n|) = i)$$

$$P(0) = \left(\frac{1}{5}\right)^n$$

$$P(\max = 2) = 1 - \left(\frac{3}{5}\right)^n$$

$$P(\max = 1) = \left(\frac{3}{5}\right)^n - \left(\frac{1}{5}\right)^n$$

$$\therefore E[\|\mathbf{x}\|_{\max}] = 1 \cdot \left(1 - \left(\frac{3}{5}\right)^n - \left(\frac{1}{5}\right)^n\right) + \left(\frac{3}{5}\right)^n \cdot 2$$

Question 1c an 2a

In [62]: Image("/content/sample_data/Screen Shot 2020-02-22 at 9.40.01 AM.png")

Out[62]:

C Since x_i 's are i.i.d.

$$\text{cov}(x_i, x_j) = \text{Var}(x_i) = E(x_i^2) - E^2(x_i) = 2 - 0 = 2$$
$$E^2(x_i) = 0$$
$$E(x_i^2) = 2$$
$$\therefore \begin{bmatrix} 2 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 2 \end{bmatrix}$$

1.2

(a) $E[(a-y)^2] = E[a^2 - 2ay + y^2]$

$$= E[a^2] + E[y^2] - 2aE[y]$$
$$= \text{Var}(a) + \text{Var}(y) + E[a^2] + E[y^2] - 2aE[y]$$
$$= \text{Var}(y) + (a - E[y])^2$$

\therefore To minimize $E(a-y)^2$; $a = E[y]$
with risk of $\text{Var}(y)$

Prove that Bayes risk = $\text{Var}(y)$

$$\begin{aligned} E(a^* - y)^2 &= E(E(y) - y)^2 \\ &= E[E(y)^2 - 2yE(y) + y^2] \\ &= E[y]^2 - 2E(y)E(y) + E[y^2] \\ &= \text{Var}(y) \end{aligned}$$

Question 2b

In [63]: Image("/content/sample_data/Screen Shot 2020-02-22 at 9.41.10 AM.png")

Out[63]:

(i) $f(x^*) = \arg \min_a E[(a-y)^2 | x]$

$$\begin{aligned} E[(a-y)^2 | x] &= E[a^2 + y^2 - 2ay | x] \\ &= E[a^2 | x] + E[y^2 | x] - E[2ay | x] \\ &= a^2 + E[y^2 | x] - 2a E[y | x] \\ \frac{\partial}{\partial a} &\Rightarrow 2a - 2E[y | x] = 0 \\ &\Rightarrow a = E[y | x] \\ \therefore f(x^*) &= E[y | x] \end{aligned}$$

(ii) $E[(f(x)-y)^2]$
 $\rightarrow E[E[(f(x)-y)^2 | x]]$

Let $g(y) = E[(f(x)-y)^2 | x]$
Given $g(y^*) = E[(f(x)-y)^2 | x]$

$$\begin{aligned} g(y^*) &\leq g(y) \\ \sum g(y^*) P(x=x) &\leq \sum g(y) P(x=x) \\ E(g(y^*)) &\leq E(g(y)) \\ \therefore E[E[(f(x)-y^*)^2 | x]] &\leq E[E[(f(x)-y)^2 | x]] \\ \therefore E[(f(x)-y^*)^2] &\leq E[(f(x)-y)^2] \end{aligned}$$

```
In [0]: import sys
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
```

```
In [0]: df = pd.read_csv('/content/sample_data/ridge_regression_dataset.csv', delimiter=',')
```

LINEAR REGRESSION:

1. Normalization :

Numpy's broadcasting can be used here. It makes the calculations easier. It handles array with various sizes and performs calculations which can be used here.

```
In [0]: def feature_normalization(train, test):
        """Rescale the data so that each feature in the training set is in
        the interval [0,1], and apply the same transformations to the test
        set, using the statistics computed on the training set.

        Args:
            train - training set, a 2D numpy array of size (num_instances, num_features)
            test - test set, a 2D numpy array of size (num_instances, num_features)

        Returns:
            train_normalized - training set after normalization
            test_normalized - test set after normalization
        """
        train_normedlist = []
        for i,v in enumerate(train):
            result = all(elem == train[i][0] for elem in train[i])
            if result == True:
                train[i] = train[i]
                train_normedlist.append(train[i])
            else:
                for k,v in enumerate(train[i]):
                    train[i][k] = (train[i][k] - min(train[i])) / (max(train[i]) - min(train[i]))
        train_normedlist.append(train[i])

        test_normedlist = []
        for i,v in enumerate(test):
            result = all(elem == test[i][0] for elem in test[i])
            if result == True:
                test[i] = test[i]
                test_normedlist.append(test[i])
            else:
                for k,v in enumerate(test[i]):
                    test[i][k] = (test[i][k] - min(train[i])) / (max(train[i]) - min(train[i]))
        test_normedlist.append(test[i])

        return np.array(train_normedlist), np.array(test_normedlist)
```

```
In [0]: X = df.values[:, :-1]
        y = df.values[:, -1]

        print('Split into Train and Test')
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =100, random_state=10)

        print("Scaling all to [0, 1]")
        X_train, X_test = feature_normalization(X_train, X_test)
        X_train = np.hstack((X_train, np.ones((X_train.shape[0], 1)))) # Add bias term
        X_test = np.hstack((X_test, np.ones((X_test.shape[0], 1)))) # Add bias term'''

        Split into Train and Test
        Scaling all to [0, 1]
```

1. a

Made Design Matrix. dmX (Design Matrix X) = X_{train} θ = θ response = y_{train}

```
In [0]: m = X_train.shape[0]

        dmX = np.zeros((m, X_train.shape[1]))
        dmX = np.array(X_train)

        theta = np.zeros((X_train.shape[1]))
```

```
In [0]: response = np.ones((m, 1))
        response = y_train
        b = np.ones((m, 1))
```

```
In [0]: dmX.shape
```

```
Out[0]: (100, 49)
```

$$J(\theta) = \frac{1}{m}(x\theta - y)^T(x\theta - y)$$

1. b

Indented block

```
gradientjtheta = (2/m)*((np.dot((response - (np.dot(dmX, theta))),dmX)))
```

$$\nabla J(\theta) = \frac{2}{m}(x\theta - y)^T x$$

1. c

```
In [0]: eta = np.random.rand()
h = np.random.rand(X_train.shape[1],1)
length = eta*np.abs(h)
thetadash = theta+ length
```

$$\theta' = \theta + length$$

$$J(\theta + \eta h) - J(\theta) = \frac{1}{m}(x\theta' - y)^T(x\theta' - y) - J(\theta)$$

1. d

$$\theta = \theta - \frac{\eta}{m}(x\theta - y)^T x$$

1. e

```
In [0]: def compute_square_loss(dmX, response, theta):
m = dmX.shape[0]
costfntheta = (1/m)*(np.dot((np.dot(dmX,theta) - response),(np.dot(dmX,theta) - response)))
return costfntheta
```

```
In [0]: compute_square_loss(dmX,response,theta)
```

```
Out[0]: 7.961518343622414
```

1. f

```
In [0]: def compute_square_loss_gradient(X, y, theta):
iter = X.shape[0]

gradientjtheta = ((2/iter)*(np.dot(((np.dot(X, theta)) - y ),X)))

#ones = np.ones(num_instances).reshape(num_instances,1)
#extended_X = np.append(X,ones,axis = 1)
'''differences = np.dot(X,theta)-y
grad = 1.0/num_instances*(np.dot(differences,X))'''
return gradientjtheta
```

```
In [0]: compute_square_loss_gradient(dmX, response, theta).shape
```

```
Out[0]: (49,)
```

1. a

```
In [0]: #See http://ufldl.stanford.edu/wiki/index.php/Gradient\_checking\_and\_advanced\_optimization
def grad_checker(X, y, theta, epsilon=0.01, tolerance=1e-4):
    """Implement Gradient Checker
    Check that the function compute_square_loss_gradient returns the
    correct gradient for the given X, y, and theta.

    Let d be the number of features. Here we numerically estimate the
    gradient by approximating the directional derivative in each of
    the d coordinate directions:
    (e_1 = (1,0,0,...,0), e_2 = (0,1,0,...,0), ..., e_d = (0,...,0,1))

    The approximation for the directional derivative of J at the point
    theta in the direction e_i is given by:
    ( J(theta + epsilon * e_i) - J(theta - epsilon * e_i) ) / (2*epsilon).

    We then look at the Euclidean distance between the gradient
    computed using this approximation and the gradient computed by
    compute_square_loss_gradient(X, y, theta). If the Euclidean
    distance exceeds tolerance, we say the gradient is incorrect.

    Args:
        X - the feature vector, 2D numpy array of size (num_instances, num_features)
        y - the label vector, 1D numpy array of size (num_instances)
        theta - the parameter vector, 1D numpy array of size (num_features)
        epsilon - the epsilon used in approximation
        tolerance - the tolerance error

    Return:
        A boolean value indicating whether the gradient is correct or not
    """
    true_gradient = compute_square_loss_gradient(X, y, theta) #The true gradient
    num_features = theta.shape[0]
    approx_grad = np.zeros(num_features)
    e=np.zeros((num_features,1))
    for i in range(num_features):
        e=np.zeros(num_features)
        e[i] = 1
        thetaplus = theta+epsilon * e
        thetaminus = theta-epsilon * e
        approx_grad[i] = (compute_square_loss(X, y, thetaplus) - compute_square_loss(X, y,
        thetaminus))/(2*epsilon)
    distance = np.linalg.norm(approx_grad-true_gradient)
    return distance<tolerance

    #thetaplusep, thetaminusep = (theta+epsilon),(theta-epsilon)

    #TODO
```

```
In [0]: grad_checker(dmX, response, theta, epsilon=0.01, tolerance=1e-4)
```

```
Out[0]: True
```

1. a

```
In [0]: def batch_grad_descent(X, y, alpha=0.1, num_step=1000, grad_check=False):
        """
        In this question you will implement batch gradient descent to
        minimize the average square loss objective.

        Args:
            X - the feature vector, 2D numpy array of size (num_instances, num_features)
            y - the label vector, 1D numpy array of size (num_instances)
            alpha - step size in gradient descent
            num_step - number of steps to run
            grad_check - a boolean value indicating whether checking the gradient when updating

        Returns:
            theta_hist - the history of parameter vector, 2D numpy array of size (num_step+1, num_features)
                           for instance, theta in step 0 should be theta_hist[0], theta in step
                           (num_step) is theta_hist[-1]
            loss_hist - the history of average square loss on the data, 1D numpy array, (num_step+1)
        """
        num_instances, num_features = X.shape[0], X.shape[1]
        theta_hist = np.zeros((num_step+1, num_features)) #Initialize theta_hist
        loss_hist = np.zeros(num_step+1) #Initialize loss_hist
        theta = np.zeros(num_features) #Initialize theta
        for i in range(1, num_step+1):
            grad = compute_square_loss_gradient(X, y, theta)
            theta = theta - alpha * grad
            loss_hist[i] = compute_square_loss(X, y, theta)
            theta_hist[i] = theta
        return theta_hist, loss_hist
```

```
In [0]: theta_hist, loss_hist = batch_grad_descent(X, y, alpha=0.1, num_step=1000, grad_check=False)
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:24: RuntimeWarning: invalid value encountered in subtract
```

1. b


```
In [0]: step_size = [0.5, 0.1, .05, .01, 1]
```



```
In [0]: for i in range(len(step_size)):
        theta_hist, loss_hist = batch_grad_descent(X, y, alpha=step_size[i], num_step=1000, grad_check=False)
        plt.plot(np.log(loss_hist), label = step_size[i])
        plt.legend()
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:24: RuntimeWarning: invalid value encountered in subtract
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:24: RuntimeWarning: invalid value encountered in subtract
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
This is separate from the ipykernel package so we can avoid doing imports until
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
This is separate from the ipykernel package so we can avoid doing imports until
```

```
Out[0]: <matplotlib.legend.Legend at 0x7f74cae0cba8>
```



Larger the steps, the faster the loss function value decreases. When step size > 0.01, it starts to diverge.

RIDGE REGRESSION!

1.

$$\nabla J(\theta) = \frac{2}{m}(x\theta - y)^T x + 2\lambda\theta$$

2.


```
In [0]: def compute_regularized_square_loss_gradient(X, y, theta, lambda_reg):
        """
        Compute the gradient of L2-regularized average square loss function given X, y and theta
        a

        Args:
            X - the feature vector, 2D numpy array of size (num_instances, num_features)
            y - the label vector, 1D numpy array of size (num_instances)
            theta - the parameter vector, 1D numpy array of size (num_features)
            lambda_reg - the regularization coefficient

        Returns:
            grad - gradient vector, 1D numpy array of size (num_features)

        """

        ridgecostfntheta = compute_square_loss_gradient(X, y, theta) + (lambda_reg*2*theta)
        return ridgecostfntheta
```

```
In [0]: ridgecostfn = compute_regularized_square_loss_gradient(dmX, response, theta, lambda_reg = 2
        )
```

3.

```
In [0]: def regularized_grad_descent(X, y, alpha=0.05, lambda_reg=10**-2, num_step=1000):
        """
        Args:
            X - the feature vector, 2D numpy array of size (num_instances, num_features)
            y - the label vector, 1D numpy array of size (num_instances)
            alpha - step size in gradient descent
            lambda_reg - the regularization coefficient
            num_step - number of steps to run

        Returns:
            theta_hist - the history of parameter vector, 2D numpy array of size (num_step+1, num_features)
                        for instance, theta in step 0 should be theta_hist[0], theta in step
                        (num_step+1) is theta_hist[-1]
            loss_hist - the history of average square loss function without the regularization
                        term, 1D numpy array.
        """

        num_instances, num_features = X.shape[0], X.shape[1]
        theta = np.zeros(num_features) #Initialize theta
        theta_hist = np.zeros((num_step+1, num_features)) #Initialize theta_hist
        loss_hist = np.zeros(num_step+1) #Initialize loss_hist
        for i in range(1, num_step+1):
            loss_hist[i] = compute_square_loss(X,y,theta)

            grad = compute_regularized_square_loss_gradient(X, y, theta, lambda_reg)
            theta = theta-alpha*grad
            theta_hist[i] = theta

        return theta_hist, loss_hist
```

```
In [0]: theta_histridge, loss_histridge = regularized_grad_descent(dmX, response, alpha=0.05, lambda_reg=10**-2, num_step=1000)
```

4.

Since B multiplied by the last theta term which is being adjusted. For larger B, the magnitude of theta adjusting is small and the regularization terms is small compared to the entire regularization. Bias is used to regularize the training set.

5.

```
In [0]: for lambda_reg in [10**-7,10**-5,10**-3,10**-1]:
        theta_hist,loss_hist= regularized_grad_descent(X_train,y_train,alpha=0.05,lambda_reg=la
mbda_reg)
        plt.plot(range(len(loss_hist)),np.log(loss_hist),label='lambda:'+str(lambda_reg))
        plt.xlabel("step")
        plt.ylabel("log(square loss)")
        plt.legend()
        plt.show()
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log

This is separate from the ipykernel package so we can avoid doing imports until

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log


This is separate from the ipykernel package so we can avoid doing imports until

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log


This is separate from the ipykernel package so we can avoid doing imports until

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log

This is separate from the ipykernel package so we can avoid doing imports until



```
In [0]: loss_train = []
        loss_test = []
        lambda_reg = [10**-7,10**-5,10**-3,0.01]
        for i in lambda_reg:
            theta_hist,loss_hist= regularized_grad_descent(dmX,response,lambda_reg=i)
            newtheta = theta_hist[-1]
            loss_hist_train = compute_square_loss(dmX,response,newtheta)
            loss_train.append(loss_hist_train)
            loss_hist_test= compute_square_loss(X_test,y_test,newtheta)
            loss_test.append(loss_hist_test)
        plt.plot(np.log(lambda_reg),np.log(loss_train), label = 'train')
        plt.plot(np.log(lambda_reg),np.log(loss_test), label = 'test')
        plt.legend()
        plt.show()
```



I would select theta with value where lambda is 1e-5. It converges the fastest.

Stochastic Gradient Descent

1.

$$f(\theta) = \frac{1}{m} \sum \nabla f_m \theta + \lambda \theta^T \theta$$

2.

$$\begin{aligned} \nabla f(\theta) &= \nabla (x\theta - y)^T (x\theta - y) + \lambda \theta^T \theta \\ E[\nabla f(\theta)] &= \frac{1}{m} \sum \nabla f_m \theta \\ \text{Therefore, } E[\nabla f(\theta)] &= \nabla J(\theta) \end{aligned}$$

3.

Initialized all parameters θ

$\theta' = \theta - \eta \nabla J(\theta)$ where $J(\theta)$ is calculated with a random subset of training set.

4.

```
In [0]: def stochastic_grad_descent(X, y, alpha=0.1, lambda_reg=1, num_epochs=1000, c = c):
        """
        In this question you will implement stochastic gradient descent with a regularization term

        Args:
            X - the feature vector, 2D numpy array of size (num_instances, num_features)
            y - the label vector, 1D numpy array of size (num_instances)
            alpha - string or float. step size in gradient descent
                NOTE: In SGD, it's not always a good idea to use a fixed step size. Usually
                it's set to 1/sqrt(t) or 1/t
                    if alpha is a float, then the step size in every iteration is alpha.
                    if alpha == "1/sqrt(t)", alpha = 1/sqrt(t)
                    if alpha == "1/t", alpha = 1/t
            lambda_reg - the regularization coefficient
            num_iter - number of epochs (i.e number of times) to go through the whole training set

        Returns:
            theta_hist - the history of parameter vector, 3D numpy array of size (num_iter, num_instances, num_features)
            loss_hist - the history of regularized loss function vector, 2D numpy array of size (num_iter, num_instances)
        """
        num_instances, num_features = X.shape[0], X.shape[1]
        theta = np.ones(num_features) #Initialize theta
        num_iter = num_epochs
        theta_hist = np.zeros((num_iter+1, num_features)) #Initialize theta_hist
        loss_hist = np.zeros(num_iter+1) #Initialize loss_hist
        for i in range(1, num_iter+1):
            if alpha == '1/sqrt(t)':
                step_size = c/np.sqrt((i+1.0))
            elif alpha == '1/t':
                step_size = c/(i+1.0)
            else:
                step_size = alpha

            grad = compute_regularized_square_loss_gradient(X_train, y_train, theta, lambda_reg)
            theta = theta - step_size * grad.T
            loss_i = compute_square_loss(X, y, theta) + np.dot(theta, theta) * lambda_reg
            loss_hist[i] = loss_i
            theta_hist[i] = theta
        return loss_hist, theta_hist
```

```
In [0]: h=stochastic_grad_descent(dmX, response, alpha=0.1, lambda_reg=1, num_epochs=1000)
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:34: RuntimeWarning: invalid value encountered in subtract
```

5.

```
In [0]: for step_size in [0.05,0.005, 0.01, 0.001]:
        loss_hist_SGD,theta_hist_SGD = stochastic_grad_descent(dmX,response,lambda_reg=1e-5,alpha=step_size)
        plt.plot(range(len(loss_hist_SGD)),np.log(loss_hist_SGD),label='lambda:'+str(step_size))
    plt.xlabel("step")
    plt.ylabel("log(l2 regularized square loss)")
    plt.legend()
    plt.show()
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
```

This is separate from the ipykernel package so we can avoid doing imports until

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
```


This is separate from the ipykernel package so we can avoid doing imports until

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
```


This is separate from the ipykernel package so we can avoid doing imports until

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: RuntimeWarning: divide by zero encountered in log
```


This is separate from the ipykernel package so we can avoid doing imports until



```
In [0]: for c in [0.1,0.01,0.001]:
        loss_hist_SGD,theta_hist_SGD = stochastic_grad_descent(dmX,response,lambda_reg=0.05,alpha='1/t', c= c)
        #print(loss_hist_SGD)
        plt.plot(range(len(loss_hist_SGD)),(loss_hist_SGD),label='lambda:'+str(c))
    plt.xlabel("step")
    plt.ylabel("log(l2 regularized square loss)")
    plt.legend()
    plt.show()
```



```
In [0]: for c in [0.1,0.01,0.001]:
        loss_hist_SGD,theta_hist_SGD = stochastic_grad_descent(dmX,response,lambda_reg=0.05,alp
ha='1/sqrt(t)', c= c)
        #print(loss_hist_SGD)
        plt.plot(range(len(loss_hist_SGD)),(loss_hist_SGD),label='lambda:'+str(c))
plt.xlabel("step")
plt.ylabel("log(l2 regularized sqare loss)")
plt.legend()
plt.show()
```



Gradient diverges with decreasing step-size. 0.01 ocnverges faster than 0.001!

Also, $1/\sqrt{t}$ converges faster than $1/t$

```
In [0]:
```